

# Parameter Efficient Fine Tuning:

*Blazing fast fine-tuning!*

Abbie Petulante  
DSI Postdoctoral Fellow



Data Science Institute

**Discovery through data.**

# Training a Model

## **Training from scratch**

Start with randomized  
weights

Train them fully to fit to a  
dataset

## **Fine-Tuning**

Start with existing  
weights

Keep tuning them with  
additional or new data

# Training a Model

## Training from scratch

Start with randomized weights

Train them fully to fit to a dataset

## Fine-Tuning

Start with existing weights

Keep tuning them with additional or new data

## PEFT

Start with existing weights and **freeze them**

Train a much smaller set of *new* parameters

# Training a Model

## Training from scratch

Start with randomized weights

Train them fully to fit to a dataset

## Fine-Tuning

Start with existing weights

Keep tuning them with additional or new data

## PEFT

Start with existing weights and **freeze them**

Train a much smaller set of *new* parameters



Less compute needed

Fewer model adaptations

# PEFT (Parameter Efficient Fine-Tuning)

- Strategies to fine tune **much** less of the model
  - Typically, <few%
  - With largely the same results!
- **Core idea**: large models are typically overparameterized - critical changes for new tasks reside in smaller, lower-dimensional subspaces
- Generally, most PEFT methods are **modular** (can be added/removed)

# Common PEFT Methods

- **Prompt/Prefix Tuning**
- Adapters
- Low-Rank Adaptations (LoRA / DoRA)

# Prompt/Prefix Tuning

**Idea:** Train *virtual tokens* (a small set of learned vectors) that steer the model to a desired behavior

- “Virtual tokens” - trainable vectors that *act like* tokens but aren’t in the vocabulary
- *learned representation vectors*, not weight updates

# Prompt/Prefix Tuning

Method	What it learns	Where it's Added
Soft-Prompt Tuning	$k$ trainable embedding vectors (no text form)	Prepended to the input embedding sequence once
Prefix Tuning	$k$ learned key/value vectors per layer	Injected into each layer's attention KV-store

- Like custom “hidden context” that *doesn't need to correspond to real words*
- Changes where embedding exists in the hidden space



# Common PEFT Methods

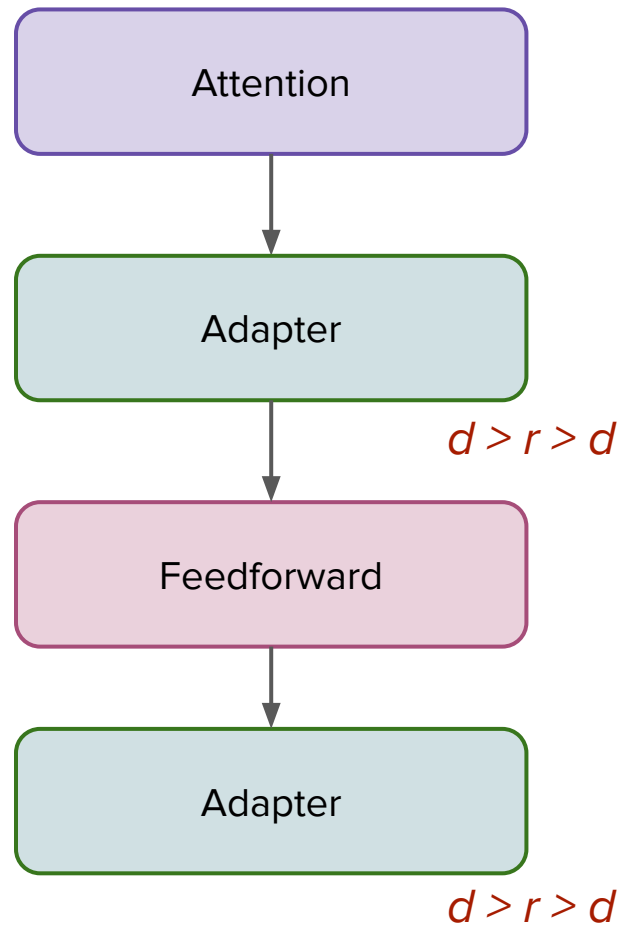

- Prompt/Prefix Tuning
- **Adapters**
- Low-Rank Adaptations (LoRA / DoRA)

# Adapters

**Idea:** Strategically insert additional “mini” layers (modules) between existing layers of the model (like plug-ins)

- Mini layers have few weights
- Train just the adapter modules, not the original model weights
- Lightweight and modular: Adapters can be swapped out for different tasks easily

# Adapters



# Adapters

Variant	Placement in Transformer block	Description / use-case
Houlsby (Sequential)	After both Attention and FFN (2 per block)	Classic bottleneck MLP; best quality, still <3 % params
Post-Attention (Bapna)	After Attention only	Fewer params/latency; popular in MT
IA <sup>3</sup> / Pfeiffer “vector”	Multiplicative scale vectors inside each block	Ultra-light (~0.05 %); merges away at inference
Parallel adapter	Runs alongside sub-layer, outputs summed	Keeps critical path short; good for real-time
Input/Output-only	One before first layer, one after last	Simplest plug-in; lowest quality but tiny

*(more common)*

# Adapters

```
from transformers import AutoModelForSequenceClassification, AutoTokenizer
from adapters import AutoAdapterModel

model = AutoAdapterModel.from_pretrained("bert-base-uncased")
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

model.add_adapter("sst2", config="houlby", reduction_factor=16)
model.train_adapter("sst2")

#set up standard training loop, train config, etc
trainer.train()

model.save_adapter("./sst2_adapter", "sst2")
```

# Common PEFT Methods

- Prompt/Prefix Tuning
- Adapters
- **Low-Rank Adaptations (LoRA + flavors)**


# Low Rank Adaptations

**Idea:** Instead of updating a whole weight matrix, add a weight update that is created from much smaller, trainable matrices

- Low rank matrices = very few parameters
  - Very quick to train!
  - Original weight matrices remain frozen

# Full Fine-Tuning

- All weights of the model are fair game to edit




$$\begin{bmatrix} n_1 \\ n_2 \\ n_3 \end{bmatrix} \times \begin{bmatrix} W_{1,4} & W_{2,4} & W_{3,4} \\ W_{1,5} & W_{2,5} & W_{3,5} \end{bmatrix} = \begin{bmatrix} n_4 \\ n_5 \end{bmatrix}$$

**All weights update, at  
each step, by  $\alpha$  \* update**



# Low Rank Adaptations (LoRA)

- Only low rank matrices A/B are tuned



$$\begin{matrix} n_1 \\ n_2 \\ n_3 \end{matrix} \times \begin{pmatrix} W_{1,4} & W_{2,4} & W_{3,4} \\ W_{1,5} & W_{2,5} & W_{3,5} \end{pmatrix} + \begin{matrix} \begin{bmatrix} \mathbf{B} \end{bmatrix} \\ d \times r \end{matrix} \begin{matrix} \begin{bmatrix} \mathbf{A} \end{bmatrix} \\ r \times d \end{matrix} = \begin{pmatrix} n_4 \\ n_5 \end{pmatrix}$$

Low-rank ( $r$ ) matrices

# Low-Rank Adaptations

LoRA Flavor	How Implemented	Extra Parameters
Standard LoRA (orig. Hu + al. 2021)	Add rank- $r$ (8–32) $\Delta W$ to chosen weight matrices	$2 \cdot d \cdot r$ per target weight
QLoRA (Dettmers 23)	Freeze 4-bit quantized base weights; train LoRA in 16-bit	Same $\Delta W$ params as standard, but less VRAM
Selective-Layer LoRA	Apply LoRA only to top-N layers <i>or</i> only to Q,V matrices	Proportional to layers chosen—often < 0.05 % total
LoRA++ / Dropout-LoRA	Adds dropout & rescaling in the LoRA branch, sometimes adapter-bias	Same as standard
DoRA (Weight-Decomposed LoRA)	Decompose each weight into magnitude ( $\gamma$ ) + direction; train $\gamma$ (scalar per row/col) and a LoRA update for only direction	Same as standard + $\gamma$ adds $O(d)$ scalars per matrix